

AGU VICTOR CHIEDOZIE

Cloud & AI Systems Engineer

WAT (GMT+1) | aguchiedoxie@gmail.com | +234 816 280 2162 | linkedin.com/in/aguchiedoxie | github.com/Moorejoe

PROFESSIONAL SUMMARY

I build cloud infrastructure and AI systems, alone or with a team, fast — usually in days rather than months. Much of what I know came from deep research and from shipping real things end-to-end, and I'm actively learning from stronger engineers every day. I'm applying for intern, junior and intermediate roles where I can ship real work and keep growing.

KEY SKILLS

AWS, Azure, Oracle Cloud (OCI), Google Cloud, Linux (Ubuntu), Nginx, Docker, Kubernetes, Terraform, GitHub Actions CI/CD, PM2 | LangChain, multi-agent orchestration, trigger-based automation design, structured function-calling, multi-tier LLM failover routing, RAG and vector agent memory | zero-trust network design, Cloudflare WAF, DDoS mitigation, DNS routing, encrypted service-to-service communication, self-healing / auto-recovery infrastructure | Python (FastAPI, Asyncio), Node.js, PostgreSQL, RESTful APIs, n8n workflow automation | Aptos MoveVM, Move smart contracts, non-custodial escrow design

TECHNICAL PROJECTS & EXPERIENCE

Designer & Builder | ForChi — AI Workflow Agent (live 24/7 Telegram bot) | REMOTE

2026 – Present

- Built and shipped a trigger-only Node.js/LangChain agent that posts to Facebook and LinkedIn automatically, using a two-step gate-and-extraction pipeline so casual chat never accidentally fires a workflow
- Designed a multi-tier LLM fallback (Gemini key rotation with a self-hosted Qwen model as backup) to stay responsive under free-tier rate limits
- Added voice-note handling and a multi-stage AI image generation pipeline, styled differently per platform
- Built an automated YouTube Shorts pipeline (AI poem → cloned voice → stock clips → assembly → upload → playlist) posting ~5 Shorts/day with topic rotation across 3 pillars: romance, philosophy, and faith
- Set up OAuth-based YouTube auth with an automatic weekly re-auth watcher, plus human-friendly self-healing error reports delivered by the bot itself

Founder & Lead Systems Architect | CloudVoid — non-custodial escrow platform (cloudvoid.online) | REMOTE

2025 – Present

- Designed and built a non-custodial P2P crypto-to-fiat escrow platform: React frontend, Python/Node.js backend, on a zero-trust Kubernetes cluster on Oracle Cloud
- Implemented an Aptos MoveVM escrow layer with pre-commit locking, trust-bond penalty logic, and gas abstraction via Paymaster sponsorship
- Built a cryptographic 3-of-5 threshold arbitration system for on-chain dispute resolution, backend locked behind Cloudflare WAF and custom iptables rules
- Set up self-healing infrastructure: a webhook-triggered watchdog tied to PM2 for automatic recovery and zero-downtime deploys

Solo Architect & Builder | Project CLAY — 21-container multi-agent system | REMOTE

2026

- Designed and built a 21-container multi-agent architecture with working image generation, conversational handling, and automated social posting
- Evaluated honestly and retired the architecture in favor of a leaner, trigger-based design that directly informed ForChi

Systems Developer | Flamchi / Footchristo — sports prediction engine (live 24/7 Telegram bot) | REMOTE

2024 – 2026

- Built a sports prediction engine for football, basketball, and tennis using association-rule mining, discarding six earlier approaches that failed on unseen data
- Every rule validated on a blind, never-touched data window: 69.6% football, 66.8% NBA, 69.8% tennis; refuses to predict when no validated pattern matches
- Latest 'Odonata dragonfly' ruleset: 950 football / 1,927 basketball / 1,495 tennis blind-validated rules
- Ships one weekly digest per sport (time · teams · projected winner) and explains WHY a pick was made via a Qwen-powered chat bot

Builder — Self-Hosting & Automation | Infrastructure (Contabo VPS, systemd, Hugging Face) | REMOTE

2026

- Migrated all production services from Render to a self-managed Contabo VPS: systemd units with auto-restart, scheduled timers, and a Telegram-driven health watchdog
- Self-hosted a local Qwen3 4B LLM (llama.cpp) on the VPS as the always-on chat fallback, cutting cold-start latency from minutes to seconds
- Distributed 148MB of media assets via a private Hugging Face dataset with a pull-on-deploy bundle so any host boots the full pipeline

CERTIFICATIONS

- Cloud Computing Certification, Digital Witch (2026)
- Certified Customer Service Professional, CURSA (2025)
- IT Support Specialist, Digital Witch (2024)